

Sysvinit 项目分析报告

李明 <limingth@gmail.com>

2013.6.22

Contents

1 Sysvinit 项目工具简介	5
1.1 项目背景介绍	5
1.2 项目架构设计	5
2 Sysvinit 项目概要分析	7
2.1 工具安装使用流程	7
2.1.1 init 命令	8
2.1.2 halt 命令	9
2.1.3 last 命令	10
2.1.4 mesg 命令	10
2.1.5 mountpoint 命令	10
2.1.6 utmpdump 命令	11
2.1.7 runlevel 命令	11
2.1.8 sulogin 命令	11
2.1.9 wall 命令	11
2.1.10 shutdown 命令	11
2.1.11 bootlogd 命令	11
2.2 代码实现概要分析	11
2.2.1 源码目录结构	11
2.2.2 Makefile 分析	13
3 Sysvinit 项目详细分析	15
3.1 init 进程代码分析	15
3.2 相关其他进程分析	15
4 Sysvinit 项目安全漏洞	17

5 Sysvinit 项目运行时调试图	19
5.1 编译安装运行调试图	19
5.1.1 wget 下载源码包	19
5.1.2 tar 解压源码包	20
5.1.3 编译项目源码	23
5.1.4 修改 Makefile 使之能够编译通过	27
5.1.5 继续编译项目源码，成功	27
5.1.6 查看生成的可执行文件	29
5.2 Linux 内核启动 init 进程	31
5.2.1 start_kernel	31
5.2.2 parse_options	31
5.2.3 rest_init	33
5.2.4 init 函数	35

Chapter 1

Sysvinit 项目工具简介

1.1 项目背景介绍

安装的程序 `halt`, `init`, `killall5`, `last`, `lastb` (链接到 `last`), `mesg`, `pidof` (链接到 `killall5`), `poweroff` (链接到 `halt`), `reboot` (链接到 `halt`), `runlevel`, `shutdown`, `sulogin`, `telinit` (链接到 `init`), `utmpdump`, `wall` 简要描述

1.2 项目架构设计

Chapter 2

Sysvinit 项目概要分析

2.1 工具安装使用流程

`halt` 正常情况下等效于 `shutdown` 加上 `-h` 参数（当前系统运行级别是 0 时除外）。它将告诉内核去中止系统，并在系统正在关闭的过程中将日志记录到 `/var/log/wtmp` 文件里。`init` 当内核已经初始化硬件，接管引导程序，开启指令线程时，`init` 会被第一个启动。`killall5` 发送一个信号到所有进程，但那些在它自己设定级别的进程将不会被这个运行的脚本所中断。`last` 给出哪一个用户最后一次登录（或退出登录），它搜索 `/var/log/wtmp` 文件，给出系统引导、关闭、运行级别改变等信息。`lastb` 给出登录失败的尝试，并写入日志 `/var/log/btmp`。`mesg` 控制是否允许其他用户也有向系统所有用户发送信息的权限。`mountpoint` 检查给定的目录是否是一个挂载点。`pidof` 报告给定程序的 PID 号。`poweroff` 告诉内核中止系统并且关闭系统（参见 `halt`）。`reboot` 告诉内核重启系统（参见 `halt`）。`runlevel` 告前一个和当前的系统运行级别，并且将最后一些运行级别写入 `/var/run/utmp`。`shutdown` 使系统安全关闭，向所有线程发送关闭信号并且通知所有已经登录的系统用户系统即将关闭。`sulogin` 允许 `root` 登录，它通常情况下是在系统在单用户模式下运行时，由 `init` 所派生。`telinit` 告诉 `init` 将切换到那一个运行级。`utmpdump` 以一个多用户友好的方式列出已经给出的登录文件的目录。`wall` 向所有已经登录的用户写入一个信息。

`Sysvinit` 软件包包含控制启动，运行和关闭所有其他程序的工具。包括：`halt` 停止系统。通常以 `-h` 参数调用 `shutdown`，但如果已经在运行级 0 的话，它就告诉内核终止系统。在这之前，它会检查文件 `/var/log/wtmp`，看系统是否正在关闭。

`init` 是所有进程的父进程。最主要的作用是在启动过程中使用 `/etc/inittab` 文件创建进程。`/etc/inittab` 文件主要为 `init` 派生 `getty` 进程提供入口，以便用户登录，在某些系统中它也被用来控制自治 (autonomous) 进程。

`killall5` 就是 `SystemV` 的 `killall` 命令。向除自己的会话 (session) 进程之外的其它进程发出信号，所以不能杀死当前使用的 `shell`。

`last` 回溯 `/var/log/wtmp` 文件（或者 `-f` 选项指定的文件），显示自从这个文件建立以来，所有用户的登录情况。

`lastb` 显示所有失败登录企图，并记录在 `/var/log/btmp`。

`mesg` 控制其它用户对用户终端的访问。

`pidof` 找出程序的进程识别号 (pid)，输出到标准输出设备。

`poweroff` 关闭系统并切断电源。但请参看 `halt`。

`reboot` 告诉内核重启系统。但请参看 `halt`。

`runlevel` 读取系统的登录记录文件 (一般是 `/var/run/utmp`) 把以前和当前的系统运行级输出到标准输出设备。

`shutdown` 以一种安全的方式终止系统，所有正在登录的用户都会收到系统将要终止通知，并且不准新的登录。

`sulogin` 允许超级用户登陆。通常是系统进入单用户模式时调用的。

`telinit` 告诉 `init` 该进入哪个运行级。

`utmpdump` 以一种用户友好的格式向标准输出设备显示 `/var/run/utmp` 文件的内容。

`wall` 向所有有信息权限的登录用户发送消息。

2.1.1 `init` 命令

`init` 命令说明

`init` 进程是所有进程的父进程。它主要的用途是根据 `/etc/inittab` 文件，来建立进程。`inittab` 文件中通常有关于登录接口的定义，就是在每个终端产生 `getty`，使用户可以进行登录。

命令格式

```
/sbin/init [ -a ] [ -s ] [ -b ] [ -z xxx ] [ 0123456Ss ]
```

运行级别

系统启动时，可以进入不同的运行级别。每个运行级别可以设置启动不同的程序。启动的每个程序都是 `init` 的进程的子进程。运行级别有 8 个：0-6, S 和 s。运行级别 0, 1 和 6 是系统保留的。运行级别 0 用来关闭系统。运行级别 1，先关闭所有用户进程和服务，然后进入单用户模式。运行级别 6 用来重启系统。运行级别 S 和 s，会直接进入单用户模式，而不进行停止进程和服务的操作。运行级别 S 和 s 的功能是相同的。

启动过程

在 kernel 启动的最后阶段，会调用 `init.init` 会查找 `/etc/inittab` 文件内容，进入指定的运行级别。当启动一个新进程时，`init` 会先检查 `/etc/initscript` 文件是否存在。如果存在，则使用这个脚本来启动那个进程。

选项

`-s, S, single`
进入单用户模式。

`1-5 Runlevel to boot into.`
进入相应的运行级别。

`-b, emergency`
直接进入单用户 shell，不运行任何其他的启动脚本。

`-Z xxx`
`-Z` 后面的参数将被忽略。可以使用这种方法将命令行加长一点，这样可以增加在堆栈中占用的空间。

2.1.2 halt 命令

停止系统。

主要选项：

`-n` reboot 或者 halt 之前，不同步 (sync) 数据。
`-w` 仅仅往 `/var/log/wtmp` 里写一个记录，并不实际做 reboot 或者 halt 操作。
`-f` 强制 halt 或者 reboot，不等其他程序退出或者服务停止就重新启动系统。这样会造成数据丢失，建议一般不要这样做。
`-i` halt 或 reboot 前，关闭所有网络接口。
`-h` halt 或 poweroff 前，使系统中所有的硬件处于等待状态。
`-p` 在系统 halt 同时，做 poweroff 操作。即停止系统同时关闭电源。

2.1.3 last 命令

2.1.4 mesg 命令

2.1.5 mountpoint 命令

查看一个目录是否为一个挂载点：

```
[root@test ~]# df
Filesystem      1K-blocks      Used Available Use% Mounted on
/dev/hda2        9918956    8036580   1370388  86% /
/dev/hda1         99043      20891     73038  23% /boot
/dev/hda5        9612604    6545956   2578352  72% /data
tmpfs            123444         0     123444   0% /dev/shm
[root@test ~]# mountpoint /
/ is a mountpoint
[root@test ~]# mountpoint /boot
/boot is a mountpoint
[root@test ~]# mountpoint /home/
/home/ is not a mountpoint
```

而且，还可以查看某个文件系统的主/从设备号：

```
[root@test ~]# df
Filesystem      1K-blocks      Used Available Use% Mounted on
/dev/hda2        9918956    8036580   1370388  86% /
/dev/hda1         99043      20891     73038  23% /boot
/dev/hda5        9612604    6545956   2578352  72% /data
tmpfs            123444         0     123444   0% /dev/shm
[root@test ~]# mountpoint -d /
3:2
[root@test ~]# mountpoint -d /boot
3:1
```

2.1.6 utmpdump 命令

2.1.7 runlevel 命令

2.1.8 sulogin 命令

2.1.9 wall 命令

wall 说明

wall 命令用来向所有用户的终端发送一条信息。发送的信息可以作为参数在命令行给出，也可在执行 wall 命令后，从终端中输入。使用终端输入信息时，按 Ctrl-D 结束输入。wall 的信息长度的限制是 20 行。

只有超级用户有权限，给所有用户的终端发送消息。

- 用法

```
usage: wall [message]
```

- 举例

```
wall ``hello msg''
```

2.1.10 shutdown 命令

2.1.11 bootlogd 命令

2.2 代码实现概要分析

2.2.1 源码目录结构

```
$ make distclean
make -C src distclean
make[1]: Entering directory `/home/akaedu/Github/sysvinit/sysvinit-2.88dsf/src'
rm -f *.o *.bak
rm -f mountpoint init halt shutdown runlevel killall5 fstab-decode sulogin bootlogd last
make[1]: Leaving directory `/home/akaedu/Github/sysvinit/sysvinit-2.88dsf/src'
$ make clean
$ tree
.
├── contrib
│   ├── alexander.viro
│   └── notify-pam-dead.patch
```

- start-stop-daemon.c
 - start-stop-daemon.README
 - TODO
 - zefram-patches
- COPYING
- COPYRIGHT
- doc
 - bootlogd.README
 - Changelog
 - Install
 - Propaganda
 - sysvinit-2.86.lsm
- Makefile
- man
 - bootlogd.8
 - bootlogd.8.todo
 - fstab-decode.8
 - halt.8
 - init.8
 - initscript.5
 - inittab.5
 - killall5.8
 - last.1
 - lastb.1
 - mesg.1
 - mountpoint.1
 - pidof.8
 - poweroff.8
 - reboot.8
 - runlevel.8
 - shutdown.8
 - sulogin.8
 - telinit.8
 - utmpdump.1
 - wall.1
- obsolete
 - bootlogd.init
 - powerd.8
 - powerd.c
 - powerd.cfg
 - powerd.README

```

|   |   | README.RIGHT.NOW
|   |   | utmpdump.c.OLD
|   | README
|   | src
|   |   | a.out
|   |   | bootlogd.c
|   |   | dowall.c
|   |   | fstab-decode.c
|   |   | halt.c
|   |   | hddown.c
|   |   | ifdown.c
|   |   | init.c
|   |   | init.h
|   |   | initreq.h
|   |   | initscript.sample
|   |   | killall5.c
|   |   | last.c
|   |   | Makefile
|   |   | mesg.c
|   |   | mountpoint.c
|   |   | oldutmp.h
|   |   | paths.h
|   |   | reboot.h
|   |   | runlevel.c
|   |   | set.h
|   |   | shutdown.c
|   |   | sulogin.c
|   |   | utmp.c
|   |   | utmpdump.c
|   |   | wall.c

```

5 directories, 69 files

2.2.2 Makefile 分析

```

93 init:          LDLIBS += $(INITLIBS) $(STATIC)
94 init:          init.o init_utm.o
95
96 halt:          halt.o ifdown.o hddown.o utmp.o reboot.h

```

```
97
98 last:          last.o oldutmp.h
99
100 mesg:         mesg.o
101
102 mountpoint:   mountpoint.o
103
104 utmpdump:     utmpdump.o
105
106 runlevel:     runlevel.o
107
108 sulogin:      LDLIBS += $(SULOGINLIBS) $(STATIC)
109 sulogin:      sulogin.o
110
111 wall:         dwall.o wall.o
112
113 shutdown:     dwall.o shutdown.o utmp.o reboot.h
114
115 bootlogd:     LDLIBS += -lutil
116 bootlogd:     bootlogd.o
```

Chapter 3

Sysvinit 项目详细分析

3.1 `init` 进程代码分析

3.2 相关其他进程分析

Chapter 4

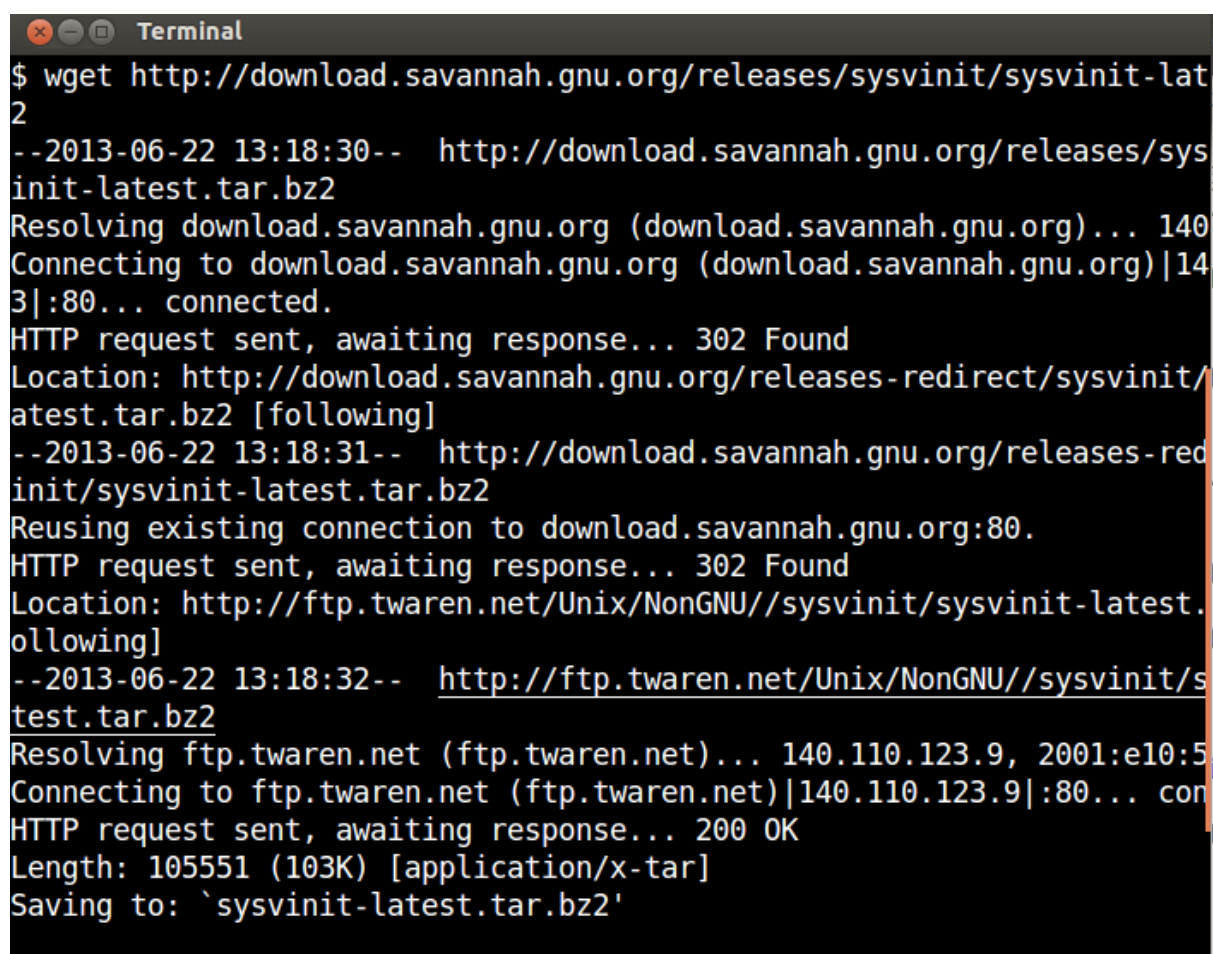
Sysvinit 项目安全漏洞

Chapter 5

Sysvinit 项目运行时调试图

5.1 编译安装运行调试图

5.1.1 wget 下载源码包



```
Terminal
$ wget http://download.savannah.gnu.org/releases/sysvinit/sysvinit-latest.tar.bz2
--2013-06-22 13:18:30-- http://download.savannah.gnu.org/releases/sysvinit-latest.tar.bz2
Resolving download.savannah.gnu.org (download.savannah.gnu.org)... 140.110.123.9
Connecting to download.savannah.gnu.org (download.savannah.gnu.org)|140.110.123.9|:80... connected.
HTTP request sent, awaiting response... 302 Found
Location: http://download.savannah.gnu.org/releases-redirect/sysvinit-latest.tar.bz2 [following]
--2013-06-22 13:18:31-- http://download.savannah.gnu.org/releases-redirect/sysvinit-latest.tar.bz2
Reusing existing connection to download.savannah.gnu.org:80.
HTTP request sent, awaiting response... 302 Found
Location: http://ftp.twaren.net/Unix/NonGNU//sysvinit/sysvinit-latest.tar.bz2 [following]
--2013-06-22 13:18:32-- http://ftp.twaren.net/Unix/NonGNU//sysvinit/sysvinit-latest.tar.bz2
Resolving ftp.twaren.net (ftp.twaren.net)... 140.110.123.9, 2001:e10:5b1:1::1
Connecting to ftp.twaren.net (ftp.twaren.net)|140.110.123.9|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 105551 (103K) [application/x-tar]
Saving to: `sysvinit-latest.tar.bz2'
```

Figure 5.1: wget 下载源码包

```

$ wget http://download.savannah.gnu.org/releases/sysvinit/sysvinit-latest.tar.bz2
--2013-06-22 13:18:30-- http://download.savannah.gnu.org/releases/sysvinit/sysvinit-l
Resolving download.savannah.gnu.org (download.savannah.gnu.org)... 140.186.70.73
Connecting to download.savannah.gnu.org (download.savannah.gnu.org)|140.186.70.73|:80.
HTTP request sent, awaiting response... 302 Found
Location: http://download.savannah.gnu.org/releases-redirect/sysvinit/sysvinit-latest
--2013-06-22 13:18:31-- http://download.savannah.gnu.org/releases-redirect/sysvinit/s
Reusing existing connection to download.savannah.gnu.org:80.
HTTP request sent, awaiting response... 302 Found
Location: http://ftp.twaren.net/Unix/NonGNU//sysvinit/sysvinit-latest.tar.bz2 [followi
--2013-06-22 13:18:32-- http://ftp.twaren.net/Unix/NonGNU//sysvinit/sysvinit-latest.t
Resolving ftp.twaren.net (ftp.twaren.net)... 140.110.123.9, 2001:e10:5c00:5::9
Connecting to ftp.twaren.net (ftp.twaren.net)|140.110.123.9|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 105551 (103K) [application/x-tar]
Saving to: `sysvinit-latest.tar.bz2'

100%[=====>] 105,551      45.1K/s   in 2.3s

2013-06-22 13:18:35 (45.1 KB/s) - `sysvinit-latest.tar.bz2' saved [105551/105551]

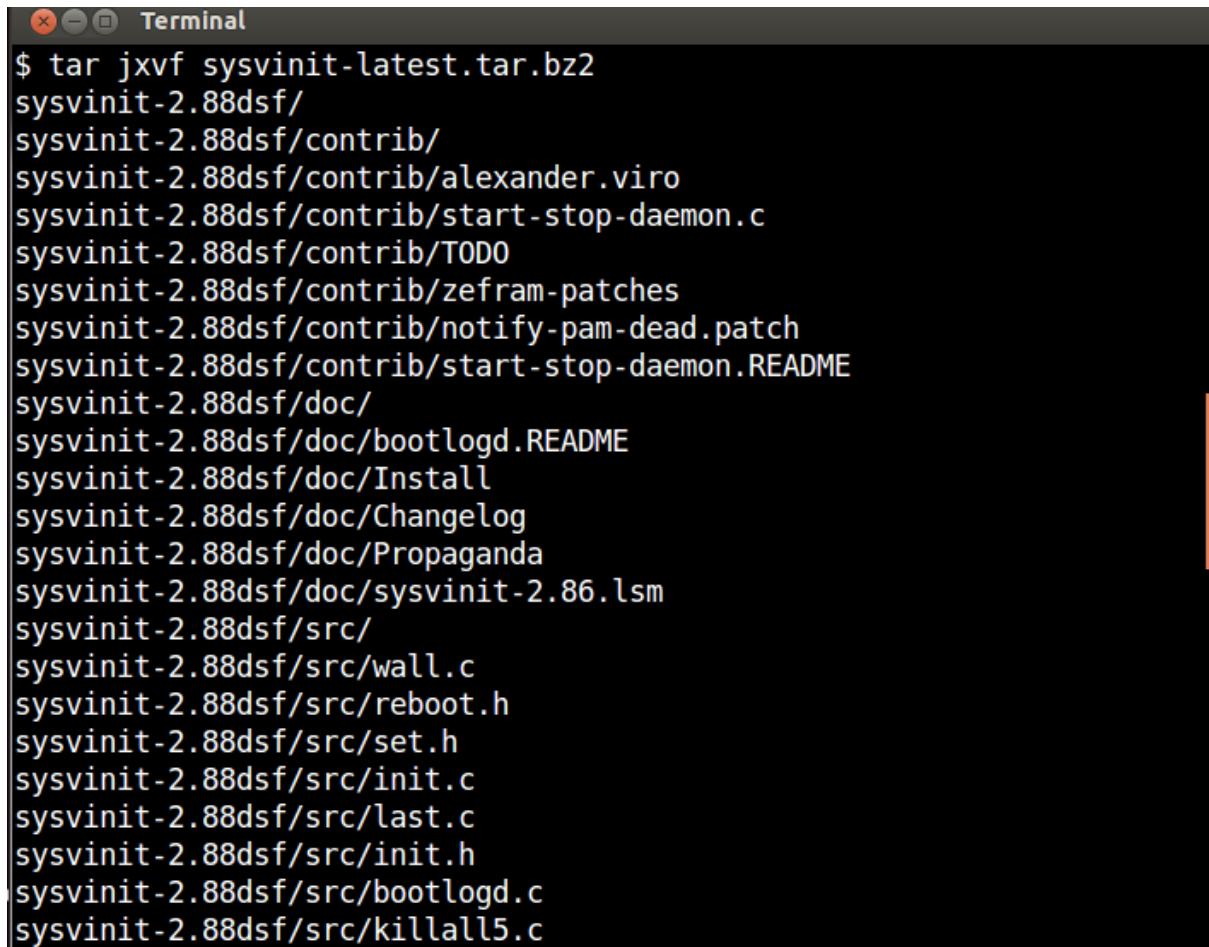
```

5.1.2 tar 解压源码包

```

$ tar jxvf sysvinit-latest.tar.bz2
sysvinit-2.88dsf/
sysvinit-2.88dsf/contrib/
sysvinit-2.88dsf/contrib/alexander.viro
sysvinit-2.88dsf/contrib/start-stop-daemon.c
sysvinit-2.88dsf/contrib/TODO
sysvinit-2.88dsf/contrib/zefram-patches
sysvinit-2.88dsf/contrib/notify-pam-dead.patch
sysvinit-2.88dsf/contrib/start-stop-daemon.README
sysvinit-2.88dsf/doc/
sysvinit-2.88dsf/doc/bootlogd.README
sysvinit-2.88dsf/doc/Install
sysvinit-2.88dsf/doc/Changelog
sysvinit-2.88dsf/doc/Propaganda
sysvinit-2.88dsf/doc/sysvinit-2.86.lsm
sysvinit-2.88dsf/src/

```

A terminal window titled "Terminal" with a dark background and light-colored text. The window shows the output of a tar command. The first line is the command prompt followed by the command. The subsequent lines are the files and directories extracted from the tar archive, listed one per line.

```
$ tar jxvf sysvinit-latest.tar.bz2
sysvinit-2.88dsf/
sysvinit-2.88dsf/contrib/
sysvinit-2.88dsf/contrib/alexander.viro
sysvinit-2.88dsf/contrib/start-stop-daemon.c
sysvinit-2.88dsf/contrib/TODO
sysvinit-2.88dsf/contrib/zefram-patches
sysvinit-2.88dsf/contrib/notify-pam-dead.patch
sysvinit-2.88dsf/contrib/start-stop-daemon.README
sysvinit-2.88dsf/doc/
sysvinit-2.88dsf/doc/bootlogd.README
sysvinit-2.88dsf/doc/Install
sysvinit-2.88dsf/doc/Changelog
sysvinit-2.88dsf/doc/Propaganda
sysvinit-2.88dsf/doc/sysvinit-2.86.lsm
sysvinit-2.88dsf/src/
sysvinit-2.88dsf/src/wall.c
sysvinit-2.88dsf/src/reboot.h
sysvinit-2.88dsf/src/set.h
sysvinit-2.88dsf/src/init.c
sysvinit-2.88dsf/src/last.c
sysvinit-2.88dsf/src/init.h
sysvinit-2.88dsf/src/bootlogd.c
sysvinit-2.88dsf/src/killall5.c
```

Figure 5.2: tar 解压源码包

```
sysvinit-2.88dsf/src/wall.c
sysvinit-2.88dsf/src/reboot.h
sysvinit-2.88dsf/src/set.h
sysvinit-2.88dsf/src/init.c
sysvinit-2.88dsf/src/last.c
sysvinit-2.88dsf/src/init.h
sysvinit-2.88dsf/src/bootlogd.c
sysvinit-2.88dsf/src/killall5.c
sysvinit-2.88dsf/src/utmpdump.c
sysvinit-2.88dsf/src/shutdown.c
sysvinit-2.88dsf/src/mountpoint.c
sysvinit-2.88dsf/src/sulogin.c
sysvinit-2.88dsf/src/fstab-decode.c
sysvinit-2.88dsf/src/initreq.h
sysvinit-2.88dsf/src/dowall.c
sysvinit-2.88dsf/src/hddown.c
sysvinit-2.88dsf/src/paths.h
sysvinit-2.88dsf/src/utmp.c
sysvinit-2.88dsf/src/ifdown.c
sysvinit-2.88dsf/src/initscript.sample
sysvinit-2.88dsf/src/halt.c
sysvinit-2.88dsf/src/oldutmp.h
sysvinit-2.88dsf/src/mesg.c
sysvinit-2.88dsf/src/Makefile
sysvinit-2.88dsf/src/runlevel.c
sysvinit-2.88dsf/COPYING
sysvinit-2.88dsf/COPYRIGHT
sysvinit-2.88dsf/man/
sysvinit-2.88dsf/man/bootlogd.8
sysvinit-2.88dsf/man/killall5.8
sysvinit-2.88dsf/man/shutdown.8
sysvinit-2.88dsf/man/bootlogd.8.todo
sysvinit-2.88dsf/man/sulogin.8
sysvinit-2.88dsf/man/fstab-decode.8
sysvinit-2.88dsf/man/mesg.1
sysvinit-2.88dsf/man/initscript.5
sysvinit-2.88dsf/man/inittab.5
sysvinit-2.88dsf/man/poweroff.8
sysvinit-2.88dsf/man/wall.1
sysvinit-2.88dsf/man/halt.8
sysvinit-2.88dsf/man/reboot.8
sysvinit-2.88dsf/man/last.1
sysvinit-2.88dsf/man/runlevel.8
sysvinit-2.88dsf/man/lastb.1
sysvinit-2.88dsf/man/pidof.8
```

```

sysvinit-2.88dsf/man/init.8
sysvinit-2.88dsf/man/utmpdump.1
sysvinit-2.88dsf/man/mountpoint.1
sysvinit-2.88dsf/man/telinit.8
sysvinit-2.88dsf/obsolete/
sysvinit-2.88dsf/obsolete/powerd.c
sysvinit-2.88dsf/obsolete/powerd.8
sysvinit-2.88dsf/obsolete/utmpdump.c.OLD
sysvinit-2.88dsf/obsolete/README.RIGHT.NOW
sysvinit-2.88dsf/obsolete/bootlogd.init
sysvinit-2.88dsf/obsolete/powerd.README
sysvinit-2.88dsf/obsolete/powerd.cfg
sysvinit-2.88dsf/Makefile
sysvinit-2.88dsf/README
$

$ ls
Makefile  pdf  sysvinit-2.88dsf  sysvinit-latest.tar.bz2

$ ls sysvinit-2.88dsf/
contrib  COPYRIGHT  Makefile  obsolete  src
COPYING  doc        man       README
$

```

5.1.3 编译项目源码

```

$ cd sysvinit-2.88dsf/
$ make
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o mountpoint.o mountpoint.
cc  mountpoint.o -o mountpoint
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o init.o init.c
init.c: In function 'telinit' :
init.c:2737:7:  warning:  ignoring return value of      'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
init.c: In function 'get_record' :
init.c:377:11:  warning:  ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:380:11:  warning:  ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:383:11:  warning:  ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]

```

```

init.c:386:11: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:389:11: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:392:11: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:395:11: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:398:11: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:401:11: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:404:11: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:423:10: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:426:10: warning: ignoring return value of      'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
init.c: In function 'spawn' :
init.c:1064:10: warning: ignoring return value of      'dup'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:1065:10: warning: ignoring return value of      'dup'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:1133:7: warning: ignoring return value of      'dup'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:1134:7: warning: ignoring return value of      'dup'
, declared with attribute warn_unused_result [-Wunused-result]
init.c: In function 'ask_runlevel' :
init.c:1673:10: warning: ignoring return value of      'write'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:1675:9: warning: ignoring return value of      'read'
, declared with attribute warn_unused_result [-Wunused-result]
init.c: In function 'make_pipe' :
init.c:1960:6: warning: ignoring return value of      'pipe'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:1965:7: warning: ignoring return value of      'write'
, declared with attribute warn_unused_result [-Wunused-result]
init.c: In function 'process_signals' :
init.c:2411:7: warning: ignoring return value of      'read'
, declared with attribute warn_unused_result [-Wunused-result]
init.c:2420:7: warning: ignoring return value of      'read'
, declared with attribute warn_unused_result [-Wunused-result]
init.c: In function 'coredump' :
init.c:666:7: warning: ignoring return value of      'chdir'

```



```

, declared with attribute warn_unused_result [-Wunused-result]
init.c: In function 'print' :
init.c:821:8: warning: ignoring return value of 'write'
, declared with attribute warn_unused_result [-Wunused-result]
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -DINIT_MAIN -c -o init_utmp.o ut
cc init.o init_utmp.o -o init
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o halt.o halt.c
halt.c: In function 'main' :
halt.c:242:2: warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o ifdown.o ifdown.c
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o hddown.o hddown.c
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o utmp.o utmp.c
cc halt.o ifdown.o hddown.o utmp.o reboot.h -o halt
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o shutdown.o shutdown.c
shutdown.c: In function 'main' :
shutdown.c:485:10: warning: variable 'realuid'
set but not used [-Wunused-but-set-variable]
shutdown.c:630:9: warning: ignoring return value of 'fscanf'
, declared with attribute warn_unused_result [-Wunused-result]
shutdown.c:719:7: warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
shutdown.c: In function 'spawn' :
shutdown.c:289:7: warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o dowall.o dowall.c
cc shutdown.o dowall.o utmp.o reboot.h -o shutdown
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o runlevel.o runlevel.c
cc runlevel.o -o runlevel
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o sulogin.o sulogin.c
sulogin.c: In function 'sushell' :
sulogin.c:407:2: warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
sulogin.c:427:8: warning: ignoring return value of 'getcwd'
, declared with attribute warn_unused_result [-Wunused-result]
cc sulogin.o -o sulogin
sulogin.o: In function `main':
sulogin.c:(.text.startup+0x1e2): undefined reference to `crypt'
collect2: ld returned 1 exit status
make: *** [sulogin] Error 1
$

```

```
Terminal
69
70 ifeq ($(WITH_SELINUX),yes)
71     SELINUX_DEF    = -DWITH_SELINUX
72     INITLIBS       += -lsepol -lselinux
73     SULOGINLIBS    = -lselinux
74 else
75     SELINUX_DEF    =
76     INITLIBS       =
77     SULOGINLIBS    =
78 endif
79
80 SULOGINLIBS       += -lcrypt
81 # Additional libs for GNU libc.
82 ifneq ($(wildcard /usr/lib*/libcrypt.a),)
83     SULOGINLIBS    += -lcrypt
84 endif
85
86 all:              $(BIN) $(SBIN) $(USBIN)
87
88 #%: %.o
89 #      $(CC) $(CFLAGS) $(LDFLAGS) -o $@ $^ $(LDLIBS)
90 #%.o: %.c
91 #      $(CC) $(CFLAGS) $(CPPFLAGS) -c $^ -o $@
"Makefile" 184L, 4343C written      80,1      42%
```

Figure 5.3: 修改 Makefile

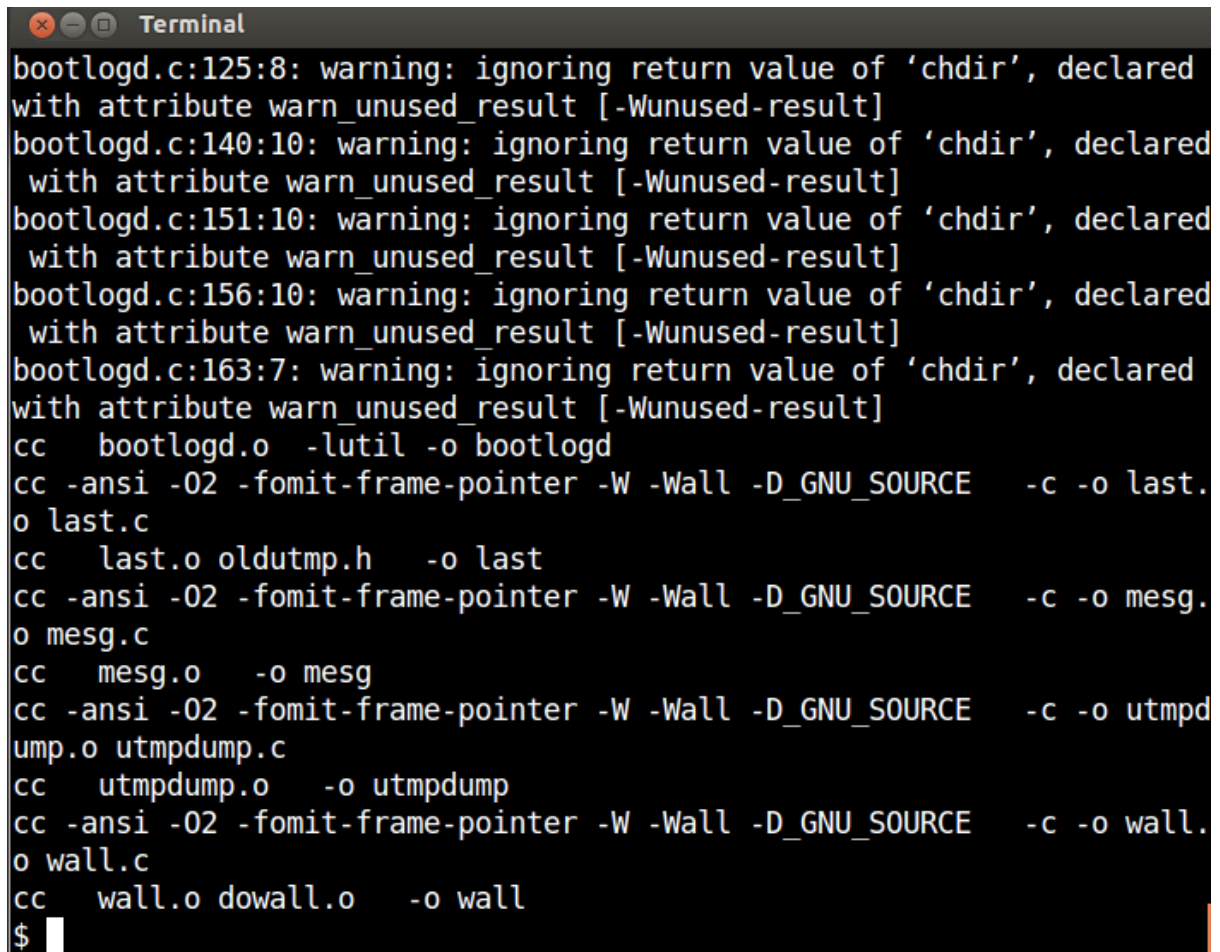
5.1.4 修改 Makefile 使之能够编译通过

```
$ vi Makefile
69
70 ifeq ($(WITH_SELINUX),yes)
71     SELINUX_DEF    = -DWITH_SELINUX
72     INITLIBS       += -lsepol -lselinux
73     SULOGINLIBS    = -lselinux
74 else
75     SELINUX_DEF    =
76     INITLIBS       =
77     SULOGINLIBS    =
78 endif
79
80 SULOGINLIBS        += -lcrypt
81 # Additional libs for GNU libc.
82 ifneq ($(wildcard /usr/lib*/libcrypt.a),)
83     SULOGINLIBS    += -lcrypt
84 endif
85
86 all:                $(BIN) $(SBIN) $(USRBIN)
87
88 #%: %.o
89 #      $(CC) $(CFLAGS) $(LDFLAGS) -o $@ $^ $(LDLIBS)
90 #%.o: %.c
91 #      $(CC) $(CFLAGS) $(CPPFLAGS) -c $^ -o $@
```

在 80 行处添加 83 行处的赋值，增加链接时 `-lcrypt` 选项

5.1.5 继续编译项目源码，成功

```
$ make
cc  sulogin.o  -lcrypt  -o sulogin
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o bootlogd.o bootlogd.c
bootlogd.c: In function 'findtty':
bootlogd.c:125:8:  warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
bootlogd.c:140:10: warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
```

A terminal window titled "Terminal" with standard window controls (close, minimize, maximize). The window contains the output of a compilation process. It shows several warning messages about ignoring the return value of 'chdir' in bootlogd.c at various line numbers. This is followed by a series of compilation commands using 'cc' to build object files (bootlogd.o, last.o, mesg.o, utmpdump.o, wall.o) and then link them into executables (last, mesg, utmpdump, wall). The prompt '\$' is visible at the bottom left.

```
bootlogd.c:125:8: warning: ignoring return value of 'chdir', declared
with attribute warn_unused_result [-Wunused-result]
bootlogd.c:140:10: warning: ignoring return value of 'chdir', declared
  with attribute warn_unused_result [-Wunused-result]
bootlogd.c:151:10: warning: ignoring return value of 'chdir', declared
  with attribute warn_unused_result [-Wunused-result]
bootlogd.c:156:10: warning: ignoring return value of 'chdir', declared
  with attribute warn_unused_result [-Wunused-result]
bootlogd.c:163:7: warning: ignoring return value of 'chdir', declared
with attribute warn_unused_result [-Wunused-result]
cc  bootlogd.o -lutil -o bootlogd
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o last.
o last.c
cc  last.o oldutmp.h -o last
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o mesg.
o mesg.c
cc  mesg.o -o mesg
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o utmpd
ump.o utmpdump.c
cc  utmpdump.o -o utmpdump
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o wall.
o wall.c
cc  wall.o dowall.o -o wall
$
```

Figure 5.4: make 编译源码包

```

bootlogd.c:151:10: warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
bootlogd.c:156:10: warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
bootlogd.c:163:7: warning: ignoring return value of 'chdir'
, declared with attribute warn_unused_result [-Wunused-result]
cc bootlogd.o -lutil -o bootlogd
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o last.o last.c
cc last.o oldutmp.h -o last
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o mesg.o mesg.c
cc mesg.o -o mesg
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o utmpdump.o utmpdump.c
cc utmpdump.o -o utmpdump
cc -ansi -O2 -fomit-frame-pointer -W -Wall -D_GNU_SOURCE -c -o wall.o wall.c
cc wall.o dowall.o -o wall
$

```

5.1.6 查看生成的可执行文件

```

$ ls -l | grep "x "
-rwxrwxr-x 1 akaedu akaedu 17677 Jun 22 13:28 a.out
-rwxrwxr-x 1 akaedu akaedu 18162 Jun 22 13:36 bootlogd
-rwxrwxr-x 1 akaedu akaedu 7402 Jun 22 13:27 fstab-decode
-rwxrwxr-x 1 akaedu akaedu 17625 Jun 22 13:30 halt
-rwxrwxr-x 1 akaedu akaedu 42121 Jun 22 13:30 init
-rwxr-xr-x 1 akaedu akaedu 706 Sep 10 2009 initscript.sample
-rwxrwxr-x 1 akaedu akaedu 22259 Jun 22 13:27 killall5
-rwxrwxr-x 1 akaedu akaedu 22117 Jun 22 13:36 last
-rwxrwxr-x 1 akaedu akaedu 7730 Jun 22 13:36 mesg
-rwxrwxr-x 1 akaedu akaedu 7708 Jun 22 13:30 mountpoint
-rwxrwxr-x 1 akaedu akaedu 7368 Jun 22 13:30 runlevel
-rwxrwxr-x 1 akaedu akaedu 27547 Jun 22 13:30 shutdown
-rwxrwxr-x 1 akaedu akaedu 17677 Jun 22 13:36 sulogin
-rwxrwxr-x 1 akaedu akaedu 12638 Jun 22 13:36 utmpdump
-rwxrwxr-x 1 akaedu akaedu 13243 Jun 22 13:36 wall
$

```

```
Terminal
$ ls -l | grep "x "
-rwxrwxr-x 1 akaedu akaedu 17677 Jun 22 13:28 a.out
-rwxrwxr-x 1 akaedu akaedu 18162 Jun 22 13:36 bootlogd
-rwxrwxr-x 1 akaedu akaedu 7402 Jun 22 13:27 fstab-decode
-rwxrwxr-x 1 akaedu akaedu 17625 Jun 22 13:30 halt
-rwxrwxr-x 1 akaedu akaedu 42121 Jun 22 13:30 init
-rwxr-xr-x 1 akaedu akaedu 706 Sep 10 2009 initscript.sample
-rwxrwxr-x 1 akaedu akaedu 22259 Jun 22 13:27 killall5
-rwxrwxr-x 1 akaedu akaedu 22117 Jun 22 13:36 last
-rwxrwxr-x 1 akaedu akaedu 7730 Jun 22 13:36 mesg
-rwxrwxr-x 1 akaedu akaedu 7708 Jun 22 13:30 mountpoint
-rwxrwxr-x 1 akaedu akaedu 7368 Jun 22 13:30 runlevel
-rwxrwxr-x 1 akaedu akaedu 27547 Jun 22 13:30 shutdown
-rwxrwxr-x 1 akaedu akaedu 17677 Jun 22 13:36 sulogin
-rwxrwxr-x 1 akaedu akaedu 12638 Jun 22 13:36 utmpdump
-rwxrwxr-x 1 akaedu akaedu 13243 Jun 22 13:36 wall
$
```

Figure 5.5: 查看可执行文件

5.2 Linux 内核启动 init 进程

5.2.1 start_kernel

```
545 asmlinkage void __init start_kernel(void)
546 {
547     char * command_line;
548     unsigned long mempages;
549     extern char saved_command_line[];
550 /*
551  * Interrupts are still disabled. Do necessary setups, then
552  * enable them
553  */
554     lock_kernel();
555     printk(linux_banner);
556     setup_arch(&command_line);
557     printk("Kernel command line: %s\n", saved_command_line);
558     parse_options(command_line);
559     trap_init();
560     init_IRQ();
561     sched_init();
562     softirq_init();
563     time_init();
564
565     .....
622     /*
623      *      We count on the initial thread going ok
624      *      Like idlers init is an unlocked kernel thread, which will
625      *      make syscalls (and thus be locked).
626      */
627     smp_init();
628     rest_init();
629 }
630
```

5.2.2 parse_options

```
426 static void __init parse_options(char *line)
427 {
428     char *next,*quote;
```

```
Terminal
545 asmlinkage void __init start_kernel(void)
546 {
547     char * command_line;
548     unsigned long mempages;
549     extern char saved_command_line[];
550 /*
551  * Interrupts are still disabled. Do necessary setups, then
552  * enable them
553  */
554     lock_kernel();
555     printk(linux_banner);
556     setup_arch(&command_line);
557     printk("Kernel command line: %s\n", saved_command_line);
558     parse_options(command_line);
559     trap_init();
560     init_IRQ();
561     sched_init();
562     softirq_init();
563     time_init();
564
565     /*
566      * HACK ALERT! This is early. We're enabling the console before
567      * we've done PCI setups etc, and console_init() must be aware of
567,17-24 66%
```

Figure 5.6: 内核 start_kernel 函数


```

429         int args, envs;
430
431         if (!*line)
432             return;
433         args = 0;
434         envs = 1;          /* TERM is set to 'linux' by default */
435         next = line;
436         while ((line = next) != NULL) {
437             quote = strchr(line, '"');
438             next = strchr(line, ' ');
439             while (next != NULL && quote != NULL && quote < next) {
440                 /* we found a left quote before the next blank
441                  * now we have to find the matching right quote
442                  */
443                 next = strchr(quote+1, '"');
444                 if (next != NULL) {
445                     quote = strchr(next+1, '"');
446                     next = strchr(next+1, ' ');
447                 }
448             }
449             if (next != NULL)
450                 *next++ = 0;
451             if (!strncmp(line, "init=", 5)) {
452                 line += 5;
453                 execute_command = line;
454             } /* In case LIL0 is going to boot us with default com

```

5.2.3 rest_init

```

532
533 static void rest_init(void)
534 {
535     kernel_thread(init, NULL, CLONE_FS | CLONE_FILES | CLONE_SIGNAL);
536     unlock_kernel();
537     current->need_resched = 1;
538     cpu_idle();
539 }
540

```

```
Terminal
425  */
426 static void __init parse_options(char *line)
427 {
428     char *next,*quote;
429     int args, envs;
430
431     if (!*line)
432         return;
433     args = 0;
434     envs = 1;      /* TERM is set to 'linux' by default */
435     next = line;
436     while ((line = next) != NULL) {
437         quote = strchr(line, '"');
438         next = strchr(line, ' ');
439         while (next != NULL && quote != NULL && quote < next) {
440             /* we found a left quote before the next blank
441              * now we have to find the matching right quote
442              */
443             next = strchr(quote+1, '"');
444             if (next != NULL) {
445                 quote = strchr(next+1, '"');
446                 next = strchr(next+1, ' ');
447             }
448         }
449         if (next != NULL)
450             *next++ = 0;
451         if (!strncmp(line,"init=",5)) {
452             line += 5;
453             execute_command = line;
454             /* In case LILO is going to boot us with default com
mand line,
@
425,1 52%
```

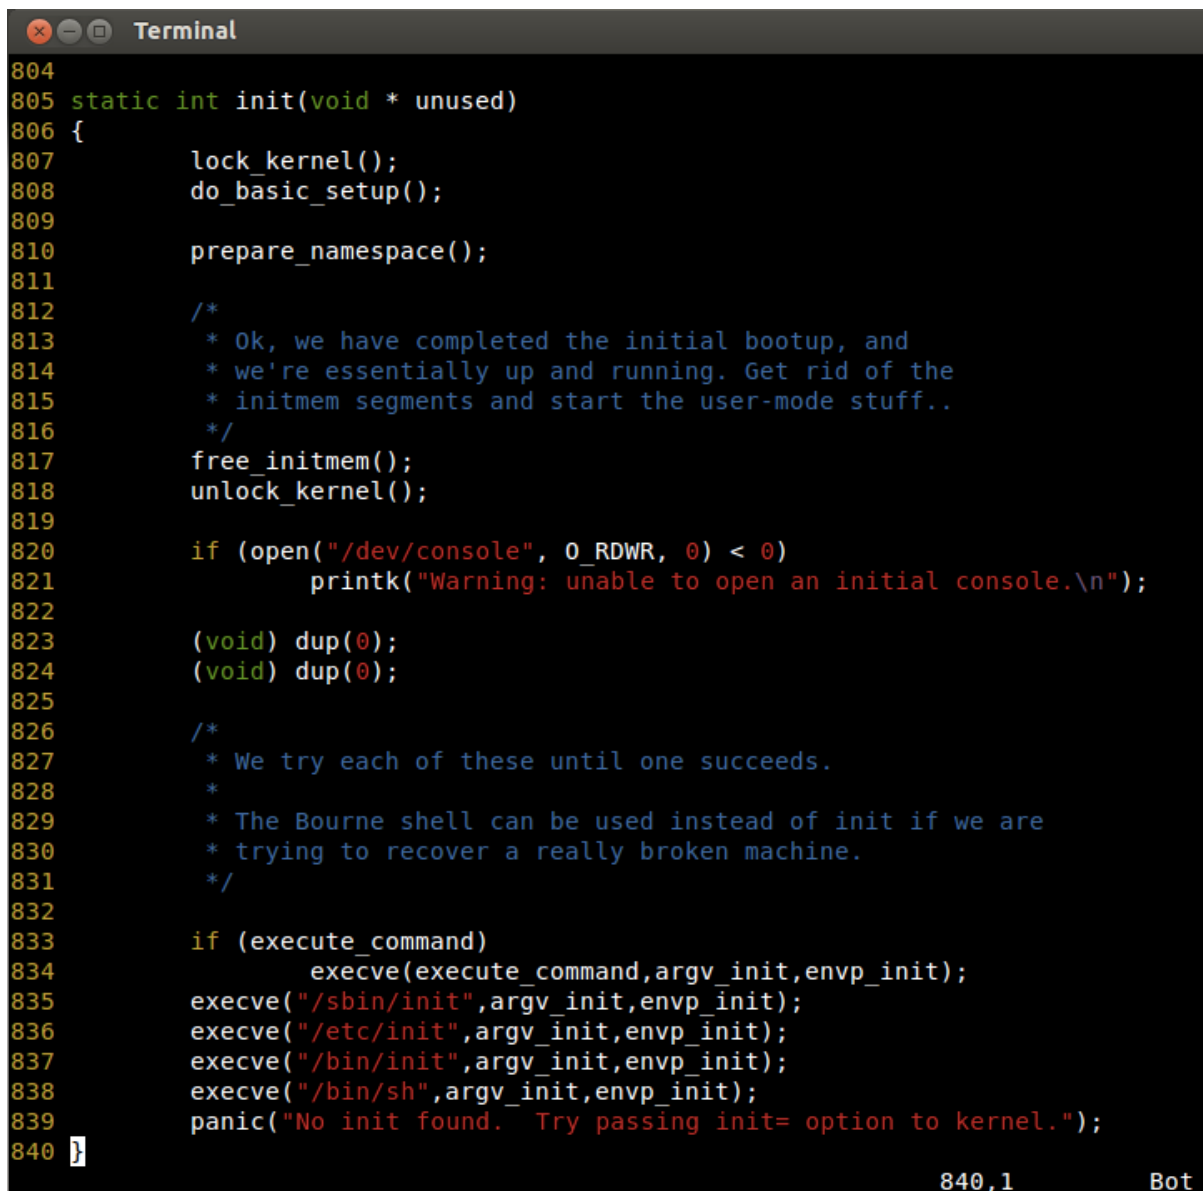
Figure 5.7: 内核 parse_options 函数

5.2.4 init 函数

```
805 static int init(void * unused)
806 {
807     lock_kernel();
808     do_basic_setup();
809
810     prepare_namespace();
811
812     /*
813      * Ok, we have completed the initial bootup, and
814      * we're essentially up and running. Get rid of the
815      * initmem segments and start the user-mode stuff..
816      */
817     free_initmem();
818     unlock_kernel();
819
820     if (open("/dev/console", O_RDWR, 0) < 0)
821         printk("Warning: unable to open an initial console.\n");
822
823     (void) dup(0);
824     (void) dup(0);
825
826     /*
827      * We try each of these until one succeeds.
828      *
829      * The Bourne shell can be used instead of init if we are
830      * trying to recover a really broken machine.
831      */
832
833     if (execute_command)
834         execve(execute_command, argv_init, envp_init);
835     execve("/sbin/init", argv_init, envp_init);
836     execve("/etc/init", argv_init, envp_init);
837     execve("/bin/init", argv_init, envp_init);
838     execve("/bin/sh", argv_init, envp_init);
839     panic("No init found. Try passing init= option to kernel.");
840 }
```

至此我们找到了一条路径，使得内核从 `start_kernel` 的主函数，进入到 `init` 进程。这里涉及到了 4 个重要的函数和 1 个重要的变量，这些都是和 `init` 进程如何启动直接相关的，对于我们了解在 `init` 进程启动之前的逻辑流程有重要作用。

- `start_kernel()`

A terminal window titled "Terminal" with a dark background and light-colored text. It displays the source code for the kernel's `init` function, with line numbers 804 through 840 on the left margin. The code includes comments in blue and function calls in green. At the bottom right of the terminal, the text "840,1" and "Bot" are visible.

```
804
805 static int init(void * unused)
806 {
807     lock_kernel();
808     do_basic_setup();
809
810     prepare_namespace();
811
812     /*
813     * Ok, we have completed the initial bootup, and
814     * we're essentially up and running. Get rid of the
815     * initmem segments and start the user-mode stuff..
816     */
817     free_initmem();
818     unlock_kernel();
819
820     if (open("/dev/console", O_RDWR, 0) < 0)
821         printk("Warning: unable to open an initial console.\n");
822
823     (void) dup(0);
824     (void) dup(0);
825
826     /*
827     * We try each of these until one succeeds.
828     *
829     * The Bourne shell can be used instead of init if we are
830     * trying to recover a really broken machine.
831     */
832
833     if (execute_command)
834         execve(execute_command, argv_init, envp_init);
835     execve("/sbin/init", argv_init, envp_init);
836     execve("/etc/init", argv_init, envp_init);
837     execve("/bin/init", argv_init, envp_init);
838     execve("/bin/sh", argv_init, envp_init);
839     panic("No init found. Try passing init= option to kernel.");
840 }
```

Figure 5.8: 内核 `init` 函数

- `parse_options()`
- `rest_init()`
- `init()`
- `execute_command`

我们用下面这张图来表示这些函数和变量之间的关系，可以更直观的看到内核启动 `init` 进程的流程。

`/sbin/init`

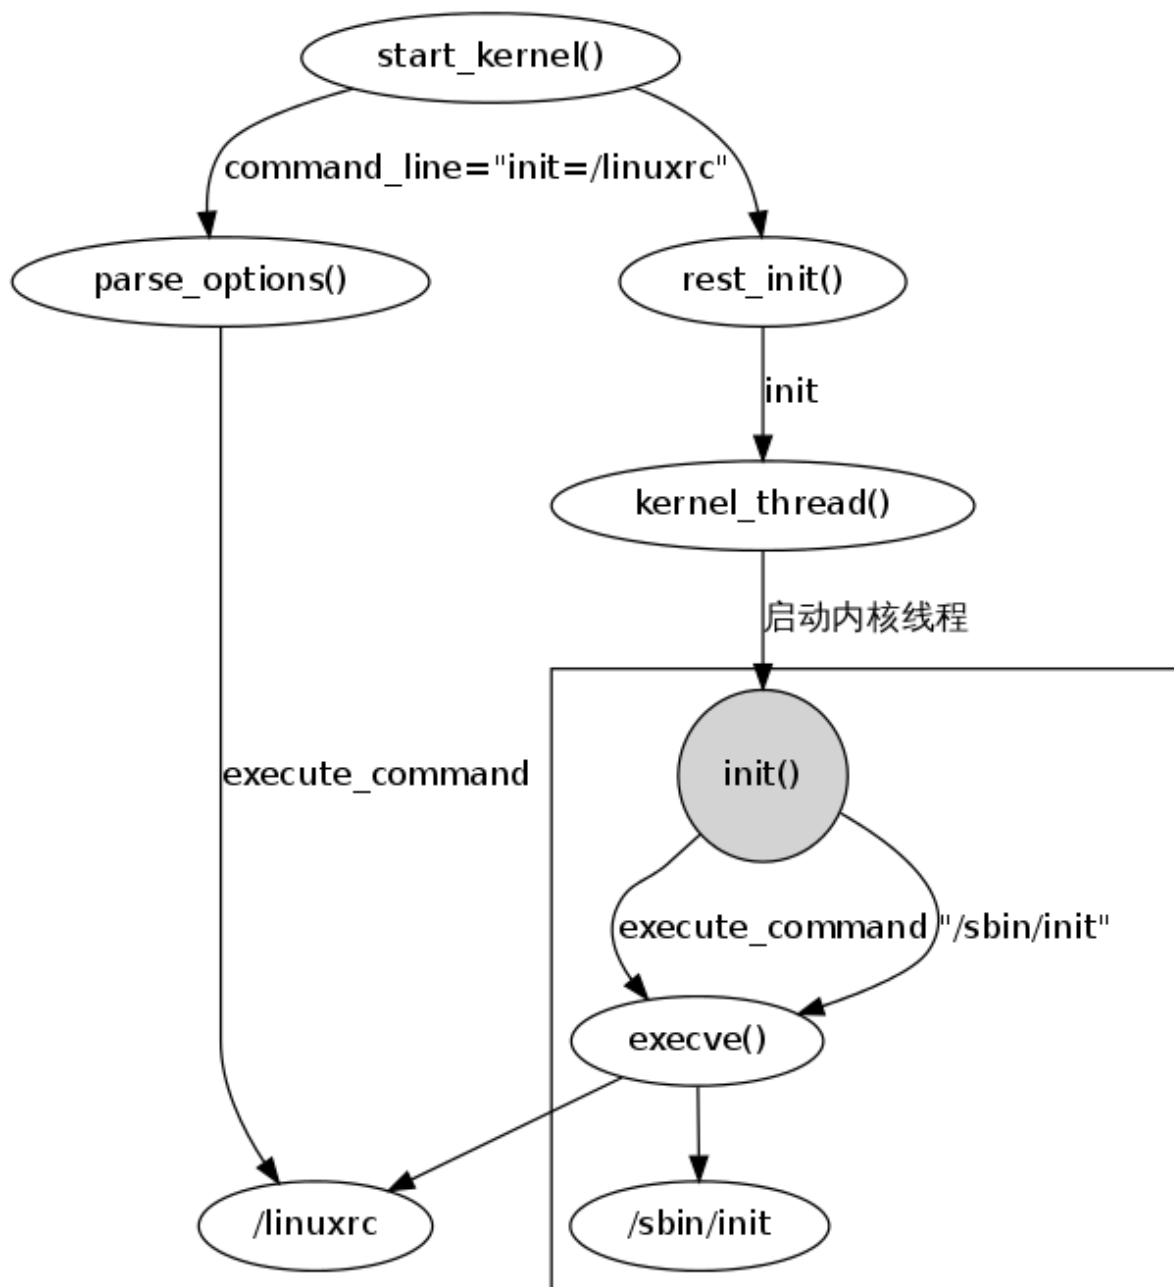


Figure 5.9: Linux 内核启动 `init` 进程